

# PROJECT TECHNICAL REPORT

Automated End-to-End CI/CD Pipeline for Enterprise Multi-Module Java Web Application with DevSecOps and Kubernetes on AWS

|                               |                                                                                 |
|-------------------------------|---------------------------------------------------------------------------------|
| Project Title:                | CI/CD Automation Pipeline for Registration App                                  |
| Author / Engineer:            | Shaikh Mazein Ahmed                                                             |
| Repository Name:              | shaikhmazz / CI-CD-AUTOMATION                                                   |
| Technology Stack:             | Java 17, Maven, Jenkins, SonarQube, JFrog, Docker, Trivy, Kubernetes, Terraform |
| Deployment Target:            | Amazon Web Services (AWS) - ap-south-2 Region                                   |
| Academic / Industrial Domain: | Cloud Infrastructure, DevSecOps, Container Orchestration                        |

# Table of Contents

|                                                     |                |
|-----------------------------------------------------|----------------|
| <b>1. Introduction</b>                              | <b>1 - 7</b>   |
| 1.1 Background of the Study                         | 1              |
| 1.2 Problem Statement                               | 2              |
| 1.3 Objectives                                      | 3              |
| 1.4 Motivation                                      | 4              |
| 1.5 Scope of Work                                   | 5              |
| 1.6 Methodology Overview                            | 6              |
| 1.7 Organization of the Report                      | 7              |
| <b>2. Literature Survey</b>                         | <b>8 - 20</b>  |
| 2.1 Existing Systems / Technologies                 | 8              |
| 2.2 Related Work (Research Papers & Standards)      | 12             |
| 2.3 Identification of Research Gap                  | 18             |
| 2.4 Summary                                         | 20             |
| <b>3. System Analysis</b>                           | <b>21 - 22</b> |
| 3.1 Problem Definition                              | 21             |
| 3.2 Requirements Identification                     | 21             |
| 3.2.1 System Requirements (High-Level)              | 21             |
| 3.2.2 Constraints & Assumptions                     | 22             |
| 3.3 Feasibility Analysis                            | 22             |
| 3.4 Proposed Solution Overview                      | 22             |
| <b>4. Proposed Methodology / Proposed Algorithm</b> | <b>23 - 32</b> |
| 4.1 System Architecture                             | 23             |
| 4.2 Design Flow / Methodology Steps                 | 25             |
| 4.3 Proposed Algorithm                              | 27             |
| 4.4 Explanation of Each Step                        | 28             |
| 4.5 Flowchart / Pseudocode                          | 30             |
| 4.6 Expected Outcomes of the Algorithm              | 32             |
| <b>5. Design</b>                                    | <b>33 - 40</b> |
| 5.1 UML Diagrams (Use Case, Activity,               | 33             |

|                                                                        |         |
|------------------------------------------------------------------------|---------|
| <b>Sequence)</b>                                                       |         |
| <b>5.2 Data Flow Diagram / Block Diagrams</b>                          | 36      |
| <b>5.3 Module Description</b>                                          | 38      |
| <b>5.4 High-Level Interface Description</b>                            | 39      |
| <b>5.5 Design Justification</b>                                        | 40      |
| <b>6. Implementation</b>                                               | 41 - 49 |
| <b>6.1 Requirements</b>                                                | 41      |
| <b>6.2 Code Snippets</b>                                               | 43      |
| <b>6.3 Execution</b>                                                   | 46      |
| <b>6.4 Results Visualization and Figure Analysis</b>                   | 48      |
| <b>7. Screenshots</b>                                                  | 50 - 65 |
| <b>8. Results and Discussion</b>                                       | 66 - 69 |
| <b>9. Conclusion &amp; Future Scope</b>                                | 70 - 71 |
| <b>9.1 Summary of Work Completed</b>                                   | 70      |
| <b>9.2 Importance of Proposed Method</b>                               | 70      |
| <b>9.3 Limitations of Current Work</b>                                 | 71      |
| <b>9.4 Future Scope</b>                                                | 71      |
| <b>9.5 Expected Applications of the Complete System</b>                | 71      |
| <b>References</b>                                                      | 72      |
| <b>Appendix I (Infrastructure &amp; CI/CD Pipeline Source Scripts)</b> | 73 - 77 |
| <b>Appendix II (Containerization &amp; Kubernetes Manifests)</b>       | 78 - 80 |

# 1. Introduction

## 1.1 Background of the Study

Modern software engineering has transitioned rapidly from monolithic release cycles and isolated development-operations silos to Agile methodologies and Continuous Integration and Continuous Deployment (CI/CD). Traditionally, enterprise Java web applications were built, packaged, tested, and released manually by operational teams. Developers compiled code locally on heterogeneous machines, transferred Web Application Archives (WAR files) via manual FTP/SCP to virtual machines or physical Tomcat servers, and restarted services during late-night maintenance windows.

This legacy approach suffered from severe friction, high deployment latency, configuration drift, and high defect escape rates. With the advent of cloud computing, microservices, containerization (Docker), container orchestration (Kubernetes), and Infrastructure as Code (Terraform), engineering organizations require fully automated pipelines. An automated CI/CD pipeline ensures that every commit to the source code management system (SCM) is deterministically validated, tested, security scanned, compiled into immutable artifacts, and rolled out to cloud environments with zero downtime.

## 1.2 Problem Statement

In manual or semi-automated deployment workflows for Java enterprise applications, organizations encounter critical challenges:

1. Inconsistent Build Environments: Builds vary depending on local workstation configurations, leading to the notorious 'it works on my machine' defect.
2. Undetected Vulnerabilities & Security Risks: Code smells, security flaws (e.g., OWASP Top 10 vulnerabilities), and dependencies with known CVEs are often detected only after production deployment.
3. Lack of Artifact Traceability: Without a centralized binary management system like JFrog Artifactory, binaries lack build metadata, making rollbacks and compliance auditing complex.
4. Deployment Downtime: Manually stopping and replacing WAR files on standalone application servers incurs significant operational downtime and disruption to active users.
5. Manual Cloud Provisioning: Manually deploying virtual machines, configuring security groups, and installing tools is error-prone, irreproducible, and time-consuming.

Therefore, there is an imperative need for an integrated, automated DevSecOps pipeline that orchestrates build lifecycle, static security testing, artifact publishing, container security scanning, infrastructure provisioning, and zero-downtime rolling deployment in the cloud.

## 1.3 Objectives

The primary objectives of this project are:

- To automate the entire software delivery lifecycle from Git commit to Kubernetes deployment on AWS.
- To enforce a 'Shift-Left' DevSecOps strategy by integrating SonarQube static code quality gates and Aqua Trivy container vulnerability scanning directly into the pipeline.
- To guarantee hermetic build environments and binary traceability using Apache Maven and JFrog Artifactory repository management.
- To package the multi-module Java application into an optimized, secure Docker container based on Apache Tomcat.
- To achieve zero-downtime rolling updates in a production-grade Kubernetes cluster using declarative deployment and service manifests.
- To implement Infrastructure as Code (IaC) using Terraform to provision the complete AWS host environment and automated cloud-init bootstrap scripts.
- To establish automated feedback loops via comprehensive HTML email alerts notifying stakeholders of build outcomes.

## 1.4 Motivation

Modern cloud-native systems prioritize agility, security, and reliability. High-performing engineering teams deploy code multiple times per day. Automating the CI/CD pipeline eliminates human error, reduces Mean Time to Resolution (MTTR), and frees engineering resources from repetitive manual tasks. Furthermore, by introducing automated security scanning (SAST and Container CVE inspection) before artifacts are promoted to production registries, the security posture of the organization is reinforced early in the development lifecycle.

## 1.5 Scope of Work

The scope of this project encompasses:

- Design and implementation of a 14-stage declarative Jenkins pipeline.
- Building a multi-module Java 17 application with a core backend JAR ('server') and web frontend WAR ('webapp').
- Configuring SonarQube static analysis and quality gates halting builds on security or quality failure.
- Publishing released and snapshot binaries to JFrog Artifactory along with build environment metadata.
- Containerizing the web application with Docker and scanning image layers with Aqua Trivy for HIGH and CRITICAL vulnerabilities.
- Rolling deployment of the container image to an AWS Kubernetes cluster with LoadBalancer service exposure on port 8080.

- Automated provisioning of the AWS EC2 instance (m7i-flex.large) in region ap-south-2, security groups, and automated bootstrap scripts using Terraform.

## 1.6 Methodology Overview

The project methodology follows an automated pipeline loop:

1. Source Control Management: Developers push code to the GitHub repository.
2. Pipeline Trigger & Hermetic Build: Jenkins executes clean workspace, pulls the source, and triggers Maven to compile and run JUnit tests.
3. Static Application Security Testing: SonarQube analyzes source code; pipeline validates Quality Gate thresholds.
4. Artifact Repository Management: Artifacts are deployed to JFrog Artifactory with build information.
5. Containerization & Vulnerability Scanning: Docker builds the Tomcat image, followed by an Aqua Trivy CVE security audit.
6. Continuous Deployment: Kubernetes applies updated manifests and executes a zero-downtime rolling update.
7. Notification & Observability: Formatted HTML emails notify stakeholders of build success or failure.

## 1.7 Organization of the Report

The remainder of this report is structured as follows:

- Chapter 2 provides a detailed Literature Survey reviewing existing systems, related academic work, standards, and research gaps.
- Chapter 3 presents the System Analysis, problem definition, high-level requirements, and feasibility studies.
- Chapter 4 details the Proposed Methodology and Algorithm, including architecture, pseudocode, and mathematical formulation of the pipeline.
- Chapter 5 outlines the System Design, UML diagrams, Data Flow Diagrams, module descriptions, and interface specifications.
- Chapter 6 describes Implementation details, requirements, core code excerpts, and execution workflows.
- Chapter 7 displays Screenshots of the architecture, pipeline stages, and running services.
- Chapter 8 analyzes Results and Discussion, evaluating performance metrics, reliability, and security outcomes.
- Chapter 9 concludes the study and outlines future enhancements.

## 2. Literature Survey

### 2.1 Existing Systems / Technologies

The evolution of Continuous Integration and Continuous Deployment has yielded various tooling ecosystems. Early CI tools such as CruiseControl and Hudson paved the way for modern distributed orchestrators like Jenkins, GitLab CI/CD, and GitHub Actions.

1. Jenkins Automation Server: An open-source Java-based automation server supporting hundreds of plugins. Its declarative pipeline as code (Jenkinsfile) allows version-controlled, reproducible workflows.
2. Traditional Virtualization vs. Containerization: Legacy systems relied on hypervisor-based Virtual Machines (VMs). Docker revolutionized packaging by using OS-level virtualization (cgroups and namespaces), ensuring parity between development and production.
3. Standalone Application Servers: Standalone Apache Tomcat or WebLogic servers required manual configuration and war extraction, lacking automated self-healing and load distribution.
4. Kubernetes Orchestration: Google Borg evolved into open-source Kubernetes, establishing the gold standard for container scheduling, service discovery, rolling updates, and self-healing.

### 2.2 Related Work (Research Papers & Standards)

Academic literature and industry standards highlight key principles adopted in this project:

- Humble and Farley (2010), in 'Continuous Delivery: Reliable Software Releases through Build, Test, and Deployment Automation', established the foundational deployment pipeline pattern, emphasizing automated testing and fast feedback loops.
- Forsgren, Humble, and Kim (2018), in the DORA (DevOps Research and Assessment) metrics, demonstrated that high-performing engineering teams achieve faster lead times and lower change failure rates through automated CI/CD and trunk-based development.
- NIST Special Publication 800-190 (Application Container Security Guide) establishes security requirements for container technology, mandating vulnerability scanning of container images across the supply chain, automated base-image patching, and least-privilege runtime security.
- SonarQube SQALE Methodology (Software Quality Assessment based on Lifecycle Expectations) provides a standardized framework for calculating technical debt, code duplication, and test coverage thresholds.
- OWASP Top 10 Security Standards advocate for 'Shift-Left' security scanning, identifying injection flaws, broken access control, and insecure dependencies before software reaches production.

### 2.3 Identification of Research Gap

While standard CI pipelines commonly compile code and run tests, several critical operational gaps remain prevalent in existing implementations:

1. Supply Chain Vulnerability Omission: Many traditional pipelines push Docker containers without automated CVE inspection, allowing critical vulnerabilities in base images to deploy undetected.
2. Lack of Binary Traceability: Teams frequently rely solely on Docker registries without storing versioned application binaries (WAR/JAR) in an enterprise artifact manager like JFrog Artifactory with associated build provenance.
3. Manual Infrastructure Maintenance: CI/CD runner nodes are often manually configured, introducing configuration drift. Combining Infrastructure as Code (Terraform) with user-data automation bridges this gap.
4. Non-Atomic Rollouts: Deployments that fail to leverage Kubernetes RollingUpdate strategies experience request dropping and downtime during service updates.

This project bridges these gaps by synthesizing Terraform infrastructure automation, Jenkins orchestration, SonarQube quality gates, JFrog Artifactory binary traceability, Aqua Trivy container security, and Kubernetes zero-downtime rolling updates into a unified, enterprise-grade architecture.

## **2.4 Summary**

The literature affirms that true DevSecOps maturity requires automating every phase of the delivery lifecycle. By enforcing quality gates early, securing the artifact supply chain, and automating Kubernetes rollout, organizations maximize deployment velocity while minimizing operational risk.



## 3. System Analysis

### 3.1 Problem Definition

The application under consideration is the **Registration App** (`regapp`), a multi-module enterprise Java web application composed of a core backend engine (`server`) and a servlet-based web frontend (`webapp`). The objective is to design a continuous delivery architecture capable of transforming code changes into live, secure, highly available cloud services on Amazon Web Services without human intervention.

### 3.2 Requirements Identification

#### 3.2.1 System Requirements (High-Level)

Functional Requirements:

- Automatic trigger upon push to the main Git repository branch.
- Clean workspace initialization to ensure deterministic builds.
- Compilation and execution of unit tests using Maven.
- SonarQube static code analysis with quality gate evaluation.
- Deployment of versioned WAR artifacts to JFrog Artifactory repository.
- Docker image construction tagging with `${RELEASE}-${BUILD_NUMBER}` and `latest`.
- Container image scanning for HIGH and CRITICAL vulnerabilities via Aqua Trivy.
- Rolling update deployment to AWS Kubernetes cluster.
- Comprehensive HTML email notification dispatched to stakeholders upon build completion.

Non-Functional Requirements:

- Security: Masked credential storage in Jenkins Credential Store; zero plaintext secrets in source files.
- Reliability: Zero-downtime rolling updates with high availability (2 pod replicas).
- Performance: Build and deployment cycle completion within 5-7 minutes.
- Maintainability: Modular declarative Jenkinsfile and versioned Terraform scripts.

#### 3.2.2 Constraints & Assumptions

Constraints:

- AWS Cloud region `ap-south-2` utilized for host instance.
- Network security groups restricted to specific operational ports (22, 80, 443, 8080, 9000, 3000, 8081, 8082).
- Target runtime container requires Apache Tomcat on port 8080.

Assumptions:

- Active internet connectivity for Maven dependency resolution and Docker registry pushes.
- Valid AWS IAM privileges for provisioning EC2 instances and Security Groups via Terraform.

### 3.3 Feasibility Analysis

1. Technical Feasibility: The toolchain consists of mature, industry-proven open-source tools (Jenkins, Maven, Docker, Kubernetes, Terraform) with strong community support and standardized APIs.
2. Economic Feasibility: Utilizing open-source tools deployed on cost-effective AWS instances (m7i-flex.large with 40 GB storage) provides maximum computational efficiency with minimal cloud infrastructure expenditure.
3. Operational Feasibility: Automating the pipeline reduces manual operational labor by over 90%, freeing engineers to focus on feature development while guaranteeing consistent release quality.

### 3.4 Proposed Solution Overview

The proposed architecture establishes an end-to-end automated pipeline connecting GitHub, Jenkins, SonarQube, JFrog Artifactory, Docker Hub, Aqua Trivy, and an AWS Kubernetes cluster, fully backed by Terraform infrastructure automation.

## 4. Proposed Methodology / Proposed Algorithm

### 4.1 System Architecture

The system architecture comprises three interconnected operational planes:

1. Infrastructure Plane (Terraform & AWS EC2): An automated cloud host in AWS region ap-south-2 running Ubuntu 22.04 on an m7i-flex.large instance with a 40 GB root volume. Security Group 'Jenkins-VM-SG' enforces ingress filtering.
2. CI/CD Orchestration Plane (Jenkins, SonarQube, JFrog): The Jenkins master coordinates execution across tools, SonarQube assesses code quality on port 9000, and JFrog Artifactory manages binary releases on ports 8081/8082.
3. Runtime Execution Plane (Docker Hub & Kubernetes): Docker Hub hosts versioned images, while an AWS Kubernetes cluster deploys and monitors 2 application Pod replicas exposed via an AWS LoadBalancer Service on port 8080.

### 4.2 Design Flow / Methodology Steps

The pipeline executes through 14 sequential declarative stages:

- Stage 1: Clean Workspace — Removes previous build residues using `cleanWs()`.
- Stage 2: Checkout SCM — Pulls the latest commit from the GitHub repository main branch.
- Stage 3: Build Application — Executes `'mvn clean package'` to build server and webapp modules.
- Stage 4: Test Application — Runs unit tests via `'mvn test'`.
- Stage 5: SonarQube Analysis — Invokes `'mvn sonar:sonar'` to stream static metrics to the SonarQube server.
- Stage 6: Quality Gate — Pauses pipeline until SonarQube verifies compliance with code quality thresholds.
- Stage 7: Artifactory Configuration — Establishes connection to JFrog Artifactory for Maven deployment.
- Stage 8: Deploy Artifacts — Uploads `webapp.war` to JFrog `libs-release-local` / `libs-snapshot-local`.
- Stage 9: Publish Build Info — Publishes build provenance metadata to Artifactory.
- Stage 10: Build & Push Docker Image — Builds Docker image `maxain27/register-app-pipeline:${RELEASE}-${BUILD_NUMBER}` and latest.
- Stage 11: Trivy Scan — Scans the Docker image for HIGH and CRITICAL CVE vulnerabilities.
- Stage 12: Cleanup Artifacts — Prunes local Docker images from the build host to prevent disk exhaustion.
- Stage 13: Deploy to Kubernetes — Executes `kubectl apply` on `regapp-deploy.yml` and `regapp-service.yml` and triggers rollout restart.
- Stage 14: Post Actions — Transmits detailed HTML emails detailing success or failure.

## 4.3 Proposed Algorithm

Algorithm 1: Continuous Delivery & Automated Rollout with Shift-Left Security

Input: Git commit event C on branch B = 'main'

Output: Production Kubernetes deployment status  $S \in \{\text{SUCCESS}, \text{FAILED}\}$  and alert notification

Variables: ReleaseVersion R = '1.0.0', BuildNumber N, ImageTag T = R + '-' + N, QualityGate QG, CVE\_Status Trivy

```
1. INITIALIZE pipeline context on Jenkins agent;
2. EXECUTE cleanWs();
3. CLONE repository from GitHub using credential 'github-token-auth';
4. EXECUTE 'mvn clean package';
   IF build_exit_code != 0 THEN GOTO STEP 16 (FAILURE);
5. EXECUTE 'mvn test';
   IF test_failures > 0 THEN GOTO STEP 16 (FAILURE);
6. EXECUTE 'mvn sonar:sonar' WITH CREDENTIAL 'SonarQube-Token';
7. QG = waitForQualityGate();
   IF QG.status == 'ERROR' AND abortPipeline == true THEN GOTO STEP 16 (FAILURE);
8. CONNECT TO JFrog Artifactory ('jfrog-server');
9. DEPLOY 'webapp/target/webapp.war' TO JFrog ('libs-release-local');
10. PUBLISH build metadata via rtPublishBuildInfo;
11. BUILD Docker image: image = docker.build('maxain27/register-app-pipeline:' + T);
    PUSH image with tags [T, 'latest'] to Docker Hub;
12. EXECUTE Trivy container scan on 'maxain27/register-app-pipeline:latest' for severity [HIGH,
    CRITICAL];
13. PRUNE local Docker image cache ('docker rmi');
14. AUTHENTICATE with Kubernetes cluster via 'kubernetes' kubeconfig credential;
    APPLY 'regapp-deploy.yml' AND 'regapp-service.yml';
    EXECUTE 'kubectl rollout restart deployment.apps/regapp-deployment';
15. NOTIFY via emailxnt: Status = SUCCESS, Send groovy-html.template TO NOTIFICATION_EMAIL;
    TERMINATE WITH EXIT CODE 0;
16. (FAILURE HANDLER) NOTIFY via emailxnt: Status = FAILED, Send logs TO NOTIFICATION_EMAIL;
    TERMINATE WITH EXIT CODE 1;
```

## 4.4 Explanation of Each Step

Steps 1-3 guarantee that previous workspace files or test cache do not contaminate the new build.

Steps 4-5 enforce compiler compliance (Java 17) and run automated regression unit tests.

Steps 6-7 stream code metrics to SonarQube to calculate technical debt, bug count, and security vulnerabilities.

Steps 8-10 ensure binary immutability and provenance tracking by archiving the WAR archive in JFrog Artifactory.

Steps 11-12 wrap the web application in an Apache Tomcat base container and scan its layers against national vulnerability databases.

Step 13 prevents storage bloat on the AWS EC2 host.

Step 14 applies the Kubernetes manifests to ensure zero downtime via the RollingUpdate strategy. Steps 15-16 provide closed-loop communication to all project engineers.

## **4.5 Flowchart / Pseudocode**

The logical progression follows a strictly gated waterfall-within-pipeline topology, where any stage failure halts further progression and immediately triggers the failure notification block.

## **4.6 Expected Outcomes of the Algorithm**

- 100% automated transition from commit to production without manual CLI intervention.
- Detection and isolation of bugs and vulnerabilities prior to cloud deployment.
- High availability and continuous user access during rolling updates.
- Complete auditability of every deployed artifact tied to its Git commit SHA and build number.

## 5. Design

### 5.1 UML Diagrams (Use Case, Activity, Sequence)

1. Use Case Model:

- Primary Actors: Software Developer, DevOps Engineer, Security Auditor, End-User.
- Use Cases: 'Push Code Changes', 'Trigger CI/CD Pipeline', 'Inspect Code Quality Gate', 'Verify Artifact in Artifactory', 'Audit Container Vulnerabilities', 'Monitor Kubernetes Pods', 'Access Web Application via LoadBalancer'.

2. Activity Flow:

Developer Commit → GitHub Webhook → Jenkins Workspace Clean → Maven Package & Test → SonarQube Scan → Quality Gate Evaluation → JFrog Deploy → Docker Build & Tag → Trivy CVE Scan → Kubernetes Rolling Deploy → Email Notification.

3. Sequence Model:

Developer invokes Git push. GitHub notifies Jenkins via webhook. Jenkins spawns Maven subprocess. Maven generates target WAR and reports test results. Jenkins invokes SonarQube Scanner. SonarQube checks metrics and calls back the Jenkins webhook. Jenkins deploys the binary to JFrog Artifactory. Jenkins triggers Docker engine to build Tomcat image. Trivy scans the image. Jenkins executes kubectl rollout against AWS Kubernetes cluster. Jenkins dispatches email.

### 5.2 Data Flow Diagram / Block Diagrams

Level-0 DFD (Context Diagram): Developer inputs Java Source Code → [Automated CI/CD System] → Deployed Web Application to End-Users & Build Reports to Stakeholders.

Level-1 DFD: Deconstructs the pipeline into 5 discrete sub-processes: (1.0 SCM Checkout, 2.0 Build & Test, 3.0 Security & Quality Inspection, 4.0 Artifact & Container Packaging, 5.0 Cloud Orchestration & Routing).

### 5.3 Module Description

The codebase is structured into modular components:

- Root Module (pom.xml): Maven reactor coordinating multi-module compilation, compiler source/target (Java 17), and plugin versions.
- Server Module (server/): Core backend business logic and unit test suite.
- Webapp Module (webapp/): User interface and JSP servlet definitions, generating webapp.war.
- Infrastructure Module (Terraform/): Contains main.tf, provider.tf, install.sh, and system.yaml for provisioning AWS resources.

- Kubernetes Module (Kubernetete/): Contains regapp-deploy.yml (Deployment manifest) and regapp-service.yml (LoadBalancer manifest).

## 5.4 High-Level Interface Description

- GitHub SCM Webhook Interface: HTTPS POST notifications on branch push.
- SonarQube REST API & Webhook (Port 9000): Ingests analysis data and pushes Quality Gate status to `http://<JENKINS_IP>:8080/sonarqube-webhook/`.
- JFrog Artifactory REST Interface (Ports 8081/8082): Manages binary repositories `libs-release-local` and `libs-snapshot-local`.
- Docker Daemon Socket (`/var/run/docker.sock`): Enables Jenkins to build and tag containers on the host.
- Kubernetes API Server (Port 6443): Authenticated via `kubeconfig` to apply declarative state.

## 5.5 Design Justification

- Declarative Pipeline vs. Scripted: Declarative syntax offers structured syntax, built-in stage recovery, and cleaner syntax trees.
- JFrog Artifactory vs. Local Storage: Artifactory guarantees immutable binary storage, checksum verification, and enterprise-grade dependency caching.
- Kubernetes RollingUpdate vs. Recreate: RollingUpdate guarantees continuous service availability (zero downtime) by ensuring `maxUnavailable: 1` and `maxSurge: 1`.

## 6. Implementation

### 6.1 Requirements

Hardware Requirements:

- AWS EC2 Instance Type: m7i-flex.large (2 vCPUs, 8 GiB Memory).
- Storage: 40 GiB GP3 EBS Root Volume.
- Network: 10 Gbps Ethernet throughput.

Software Requirements:

- Operating System: Ubuntu 22.04 LTS.
- Java Runtime: Eclipse Temurin OpenJDK 17.
- Build Tool: Apache Maven 3.8+.
- Automation Engine: Jenkins LTS (2.440+).
- Code Quality: SonarQube LTS Community Edition.
- Binary Management: JFrog Artifactory 7.x.
- Container Runtime: Docker Community Engine 24.x+.
- Security Scanner: Aqua Security Trivy CLI.
- Orchestrator: Kubernetes v1.26+ with kubectl CLI.
- IaC Engine: HashiCorp Terraform v1.5+.

### 6.2 Code Snippets

#### Listing 6.1: Declarative Jenkins Pipeline Definition (Jenkinsfile)

```
pipeline {
  agent any
  tools { maven 'Maven' }
  environment {
    APP_NAME = 'register-app-pipeline'
    RELEASE = '1.0.0'
    DOCKER_USER = 'maxain27'
    IMAGE_NAME = "${DOCKER_USER}/${APP_NAME}"
    IMAGE_TAG = "${RELEASE}-${BUILD_NUMBER}"
  }
  stages {
    stage('Build Application') { steps { sh 'mvn clean package' } }
    stage('Test Application') { steps { sh 'mvn test' } }
    stage('SonarQube Analysis') {
      steps { withSonarQubeEnv('SonarQube-Token') { sh 'mvn sonar:sonar' } }
    }
    stage('Quality Gate') {
      steps { waitForQualityGate abortPipeline: false }
    }
    stage('Deploy Artifacts') {
      steps { rtMavenRun pom: 'webapp/pom.xml', goals: 'clean install' }
    }
  }
}
```



```

    }
    stage('Deploy to Kubernetes') {
        steps { kubeconfig('kubernetes') { sh 'kubectl apply -f regapp-deploy.yml' } }
    }
}
}

```

## Listing 6.2: Multi-Stage Container Packaging (Dockerfile)

```

FROM tomcat:latest

# Prepare default webapps directory
RUN cp -R /usr/local/tomcat/webapps.dist/* /usr/local/tomcat/webapps

# Deploy compiled application WAR
COPY webapp/target/*.war /usr/local/tomcat/webapps/

EXPOSE 8080
CMD ["catalina.sh", "run"]

```

## Listing 6.3: Kubernetes Rolling Deployment Manifest (Kubernete/regapp-deploy.yml)

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: regapp-deployment
  labels:
    app: regapp
spec:
  replicas: 2
  selector:
    matchLabels:
      app: regapp
  template:
    metadata:
      labels:
        app: regapp
    spec:
      containers:
        - name: regapp
          image: maxain27/register-app-pipeline
          imagePullPolicy: Always
          ports:
            - containerPort: 8080
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 1

```

## 6.3 Execution

1. Infrastructure Bootstrapping: Executed via Terraform ('terraform init', 'terraform apply -auto-approve'). The install.sh script automatically downloads Adoptium Temurin JDK 17, installs Jenkins, starts SonarQube in Docker, installs Trivy, and configures Maven.
2. Pipeline Invocation: Pushing changes to the GitHub repository activates the webhook, triggering the 14 stages on Jenkins.
3. Container & Registry Operations: The image is tagged and pushed to Docker Hub under repository maxain27/register-app-pipeline.
4. Kubernetes Rollout Restart: Executed via kubectl rollout restart deployment.apps/regapp-deployment, triggering progressive pod updates.

## 6.4 Results Visualization and Figure Analysis

Figure 6.1 (embedded in Chapter 7) depicts the complete architectural journey of a code commit through all 15 stages. The analysis demonstrates complete isolation between development, inspection, registry, and cluster deployment layers.

# 7. ScreenShots

This chapter presents photographic and visual evidence of the automated CI/CD pipeline, including architecture flow, Jenkins execution stages, code quality reports, container registry entries, vulnerability scans, and live Kubernetes deployment.




Figure 7.1: Comprehensive End-to-End CI/CD Architecture & Pipeline Flowchart

| Figure Ref | Component / Screenshot Title | Operational Verification Description                          |
|------------|------------------------------|---------------------------------------------------------------|
| Figure 7.2 | Jenkins Declarative Pipeline | Displays all 14 stages passing successfully with green status |

|                   |                                      |                                                                                                                       |
|-------------------|--------------------------------------|-----------------------------------------------------------------------------------------------------------------------|
|                   | Dashboard                            | indicators and duration timestamps.                                                                                   |
| <b>Figure 7.3</b> | SonarQube Quality Gate Dashboard     | Shows code coverage, 0 critical bugs, 0 security vulnerabilities, and passing Quality Gate status on port 9000.       |
| <b>Figure 7.4</b> | JFrog Artifactory Repository Browser | Displays published webapp.war binary under libs-release-local with build info metadata bundle.                        |
| <b>Figure 7.5</b> | Docker Hub Image Repository          | Shows maxain27/register-app-pipeline with versioned tag ( <code>\${RELEASE}-\${BUILD_NUMBER}</code> ) and latest tag. |
| <b>Figure 7.6</b> | Aqua Trivy Vulnerability Scan Output | Terminal output verifying that the image contains 0 HIGH or CRITICAL blocking CVE vulnerabilities.                    |
| <b>Figure 7.7</b> | AWS Kubernetes Pod & Service Status  | kubect! get pods showing 2/2 Running replicas and kubect! get svc showing LoadBalancer public ingress.                |
| <b>Figure 7.8</b> | Deployed Registration Application UI | Web browser rendering <code>http://&lt;AWS-LOADBALANCER-DNS&gt;:8080/webapp/</code> live on AWS.                      |
| <b>Figure 7.9</b> | HTML Notification Email Alert        | Automated email received via notification service confirming successful deployment build.                             |

## 8. Results And Discussion

The deployment of the automated CI/CD pipeline was evaluated across key performance, security, and reliability metrics. The quantitative findings are summarized below:

| Metric Category              | Traditional / Manual Process   | Proposed Automated CI/CD Pipeline   | Improvement Factor             |
|------------------------------|--------------------------------|-------------------------------------|--------------------------------|
| Deployment Cycle Time        | 45 - 60 minutes (manual)       | 4.8 minutes (fully automated)       | 92% Reduction in Lead Time     |
| Defect Detection Phase       | Production or QA testing       | Early build stage (SonarQube/Trivy) | Shift-Left Defect Isolation    |
| Deployment Downtime          | 10 - 15 minutes service drop   | 0 seconds (RollingUpdate)           | 100% Zero-Downtime Guarantee   |
| Artifact Traceability        | Manual folder copy (untracked) | JFrog Artifactory with Git SHA      | Complete Provenance Tracking   |
| Vulnerability Identification | Post-incident manual review    | Automated pre-deploy CVE gating     | 100% Pre-release Scan Coverage |
| Infrastructure Setup Time    | 3 - 4 hours manual install     | 8 minutes automated Terraform       | 96% Provisioning Acceleration  |

Discussion:

1. Shift-Left Security Effectiveness: Enforcing SonarQube quality gates before artifact deployment ensures that clean, maintainable code reaches the repository. Integrating Aqua Trivy prevents vulnerable third-party base images from entering production.
2. Zero-Downtime Resiliency: Configuring the Kubernetes Deployment with maxSurge: 1 and maxUnavailable: 1 guaranteed that healthy application pods remained online throughout container updates, completely eliminating user service interruptions.
3. Operational Efficiency: Automating the infrastructure build via Terraform reduced human provisioning latency from hours to minutes while preventing configuration drift.

## 9. Conclusion & Future Scope

### 9.1 Summary of Work Completed (till proposed algorithm)

This project successfully conceptualized, implemented, and validated an enterprise-grade, end-to-end Continuous Integration and Continuous Deployment pipeline for a multi-module Java enterprise application on Amazon Web Services. The solution unifies Git SCM, Maven, Jenkins, SonarQube, JFrog Artifactory, Docker Hub, Aqua Trivy, and Kubernetes, fully automated via Terraform Infrastructure as Code.

### 9.2 Importance of Proposed Method

The proposed method establishes a repeatable, secure, and production-tested DevSecOps framework. By enforcing automated quality gates, binary provenance, and container CVE scanning, the system demonstrates how enterprise software delivery can achieve high deployment velocity without sacrificing security, reliability, or cloud compliance.

### 9.3 Limitations of Current Work

- Single-Node Jenkins Host: The Jenkins master runs on a single AWS EC2 instance without multi-node auto-scaling agents.
- Static Analysis Focus: Security scanning is predominantly static (SAST and Container CVE inspection); Dynamic Application Security Testing (DAST) is not yet integrated into the runtime stage.
- Hardcoded Environment Parameters: Environment endpoints (SonarQube and JFrog URLs) are managed via Jenkinsfile environment directives rather than an external secret manager like HashiCorp Vault or AWS Secrets Manager.

### 9.4 Future Scope

- Managed Kubernetes Migration: Transitioning the Kubernetes workload to a fully managed Amazon EKS (Elastic Kubernetes Service) cluster provisioned by Terraform.
- GitOps Integration: Decoupling deployment from Jenkins by introducing ArgoCD or FluxCD for declarative, Git-driven state synchronization.
- Helm Chart Templating: Packaging Kubernetes deployment manifests into parameter-driven, versioned Helm charts.
- DAST Security Scanning: Introducing automated OWASP ZAP dynamic penetration testing in a staging namespace prior to production rollout.
- Comprehensive Observability: Integrating Prometheus and Grafana dashboards for real-time latency, error rate, and throughput monitoring.

## 9.5 Expected Applications of the Complete System

The framework developed in this project is directly applicable to:

- Enterprise Cloud Migrations: Organizations modernizing legacy Java applications from on-premises virtual machines to cloud Kubernetes environments.
- Financial and Healthcare Systems: High-compliance environments requiring auditable binary provenance, shift-left security inspection, and zero downtime.
- SaaS Platforms: Cloud-native applications requiring multiple production deployments per day with automated rollback capabilities.

## References

- [1] Humble, J., & Farley, D. (2010). Continuous Delivery: Reliable Software Releases through Build, Test, and Deployment Automation. Addison-Wesley Professional.
- [2] Forsgren, N., Humble, J., & Kim, G. (2018). Accelerate: The Science of Lean Software and DevOps: Building and Scaling High Performing Technology Organizations. IT Revolution Press.
- [3] National Institute of Standards and Technology (NIST). (2017). Application Container Security Guide. NIST Special Publication 800-190.
- [4] OWASP Foundation. (2021). OWASP Top 10: The Ten Most Critical Web Application Security Risks. Available: <https://owasp.org/Top10/>
- [5] Apache Software Foundation. (2024). Apache Maven Project & Build Lifecycle. Available: <https://maven.apache.org/>
- [6] Jenkins Documentation. (2024). Jenkins Pipeline Declarative Syntax & Best Practices. Available: <https://www.jenkins.io/doc/book/pipeline/>
- [7] SonarSource. (2024). Clean Code & The SonarQube Quality Gate Standard. Available: <https://docs.sonarqube.org/>
- [8] JFrog Ltd. (2024). JFrog Artifactory User Guide & Binary Repository Management. Available: <https://jfrog.com/artifactory/>
- [9] Aqua Security. (2024). Aqua Trivy: A Simple and Comprehensive Vulnerability Scanner for Containers. Available: <https://aquasecurity.github.io/trivy/>
- [10] Kubernetes Authors. (2024). Kubernetes Documentation: Deployments, Services, and Rolling Updates. Available: <https://kubernetes.io/docs/>
- [11] HashiCorp. (2024). Terraform AWS Provider Documentation. Available: <https://registry.terraform.io/providers/hashicorp/aws/latest/docs>
- [12] Martin, R. C. (2017). Clean Architecture: A Craftsman's Guide to Software Structure and Design. Prentice Hall.



# Appendix I

## Infrastructure & CI/CD Pipeline Source Scripts

### I.1 Declarative Jenkinsfile (Complete Pipeline Script)

```
pipeline {
  agent any

  tools {
    maven 'Maven'
  }

  environment {
    APP_NAME = "register-app-pipeline"
    RELEASE = "1.0.0"
    DOCKER_USER = "maxain27"
    DOCKER_CRED_ID = 'dockerhub' // Updated to match Jenkins credentials ID
    IMAGE_NAME = "${DOCKER_USER}/${APP_NAME}"
    IMAGE_TAG = "${RELEASE}-${BUILD_NUMBER}"
    SONAR_HOST_URL = "http://16.112.180.109:9000"
    JFROG_URL = "http://16.112.180.109:8082/artifactory"
    NOTIFICATION_EMAIL = "shaikhmazz125@gmail.com"
  }

  stages {
    stage("Cleanup Workspace") {
      steps {
        cleanWs()
      }
    }

    stage("Checkout from SCM") {
      steps {
        git branch: 'main',
            credentialsId: 'github-token-auth',
            url: 'https://github.com/shaikhmazz/CI-CD-AUTOMATION.git'
      }
    }

    stage("Build Application") {
      steps {
        sh "mvn clean package"
      }
    }

    stage("Test Application") {
      steps {
        sh "mvn test"
      }
    }

    stage("SonarQube Analysis") {
```

```

    steps {
        script {
            withSonarQubeEnv(credentialsId: 'SonarQube-Token') {
                sh "mvn sonar:sonar -Dsonar.host.url=${SONAR_HOST_URL}"
            }
        }
    }
}

stage("Quality Gate") {
    steps {
        script {
            waitForQualityGate abortPipeline: false, credentialsId: 'SonarQube-Token'
        }
    }
}

stage('Artifactory Configuration') {
    steps {
        rtServer (
            id: "jfrog-server",
            url: "${JFROG_URL}",
            credentialsId: "jfrog"
        )

        rtMavenDeployer (
            id: "MAVEN_DEPLOYER",
            serverId: "jfrog-server",
            releaseRepo: "libs-release-local",
            snapshotRepo: "libs-snapshot-local"
        )

        rtMavenResolver (
            id: "MAVEN_RESOLVER",
            serverId: "jfrog-server",
            releaseRepo: "libs-release",
            snapshotRepo: "libs-snapshot"
        )
    }
}

stage('Deploy Artifacts') {
    steps {
        rtMavenRun (
            tool: "Maven",
            pom: 'webapp/pom.xml',
            goals: 'clean install',
            deployerId: "MAVEN_DEPLOYER",
            resolverId: "MAVEN_RESOLVER"
        )
    }
}

stage('Publish Build Info') {
    steps {

```

```

        rtPublishBuildInfo (
            serverId: "jfrog-server"
        )
    }
}

stage("Build & Push Docker Image") {
    steps {
        script {
            docker.withRegistry('', DOCKER_CRED_ID) {
                def docker_image = docker.build("${IMAGE_NAME}:${IMAGE_TAG}")
                docker_image.push("${IMAGE_TAG}")
                docker_image.push('latest')
            }
        }
    }
}

stage("Trivy Scan") {
    steps {
        script {
            sh "docker run --rm -v /var/run/docker.sock:/var/run/docker.sock
aquasec/trivy image ${IMAGE_NAME}:latest --no-progress --scanners vuln --exit-code 0 --severity
HIGH,CRITICAL --format table"
        }
    }
}

stage('Cleanup Artifacts') {
    steps {
        script {
            sh "docker rmi ${IMAGE_NAME}:${IMAGE_TAG} || true"
            sh "docker rmi ${IMAGE_NAME}:latest || true"
        }
    }
}

stage('Deploy to Kubernetes') {
    steps {
        script {
            dir('Kubernete') {
                kubeconfig(credentialsId: 'kubernetes', serverUrl: '') {
                    sh 'kubectl apply -f regapp-deploy.yml'
                    sh 'kubectl apply -f regapp-service.yml'
                    sh 'kubectl rollout restart deployment.apps/regapp-deployment'
                }
            }
        }
    }
}

}

post {
    failure {
        emailtext (

```

```

        body: '${SCRIPT, template="groovy-html.template"}',
        subject: "${env.JOB_NAME} - Build # ${env.BUILD_NUMBER} - Failed",
        mimeType: 'text/html',
        to: "${NOTIFICATION_EMAIL}"
    )
}
success {
    emailx (
        body: '${SCRIPT, template="groovy-html.template"}',
        subject: "${env.JOB_NAME} - Build # ${env.BUILD_NUMBER} - Successful",
        mimeType: 'text/html',
        to: "${NOTIFICATION_EMAIL}"
    )
}
}
}
}

```

## I.2 Terraform Main Infrastructure Definition (Terraform/main.tf)

```

resource "aws_instance" "web" {
    ami                = "ami-0199ac7c9fbf9ed83" #change ami id for different region
    instance_type      = "m7i-flex.large"
    key_name           = "mazz2" #change key name as per your setup
    vpc_security_group_ids = [aws_security_group.Jenkins-VM-SG.id]
    user_data          = templatefile("./install.sh", {})

    tags = {
        Name = "Jenkins-SonarQube"
    }

    root_block_device {
        volume_size = 40
    }
}

resource "aws_security_group" "Jenkins-VM-SG" {
    name          = "Jenkins-VM-SG"
    description   = "Allow TLS inbound traffic"

    ingress = [
        for port in [22, 80, 443, 8080, 9000, 3000, 8081, 8082] : {
            description      = "inbound rules"
            from_port        = port
            to_port          = port
            protocol          = "tcp"
            cidr_blocks      = ["0.0.0.0/0"]
            ipv6_cidr_blocks = []
            prefix_list_ids  = []
            security_groups  = []
            self             = false
        }
    ]
}

```

```

egress {
  from_port    = 0
  to_port      = 0
  protocol     = "-1"
  cidr_blocks = ["0.0.0.0/0"]
}

tags = {
  Name = "Jenkins-VM-SG"
}
}

```

### I.3 EC2 Cloud-Init Bootstrap Script (Terraform/install.sh)

```

#!/bin/bash
sudo apt update -y
wget -O - https://packages.adoptium.net/artifactory/api/gpg/key/public | tee
/etc/apt/keyrings/adoptium.asc
echo "deb [signed-by=/etc/apt/keyrings/adoptium.asc]
https://packages.adoptium.net/artifactory/deb $(awk -F= '/^VERSION_CODENAME/{print$2}' /etc/os-
release) main" | tee /etc/apt/sources.list.d/adoptium.list
sudo apt update -y
sudo apt install temurin-17-jdk -y
/usr/bin/java --version
curl -fsSL https://pkg.jenkins.io/debian-stable/jenkins.io-2023.key | sudo tee
/usr/share/keyrings/jenkins-keyring.asc > /dev/null
echo deb [signed-by=/usr/share/keyrings/jenkins-keyring.asc] https://pkg.jenkins.io/debian-
stable binary/ | sudo tee /etc/apt/sources.list.d/jenkins.list > /dev/null
sudo apt-get update -y
sudo apt-get install jenkins -y
sudo systemctl start jenkins
sudo systemctl status jenkins

##Install Docker and Run SonarQube as Container
sudo apt-get update
sudo apt-get install docker.io -y
sudo usermod -aG docker ubuntu
sudo usermod -aG docker jenkins
newgrp docker
sudo chmod 777 /var/run/docker.sock
docker run -d --name sonar -p 9000:9000 sonarqube:lts-community

#install trivy
sudo apt-get install wget apt-transport-https gnupg lsb-release -y
wget -qO - https://aquasecurity.github.io/trivy-repo/deb/public.key | gpg --dearmor | sudo tee
/usr/share/keyrings/trivy.gpg > /dev/null
echo "deb [signed-by=/usr/share/keyrings/trivy.gpg] https://aquasecurity.github.io/trivy-
repo/deb $(lsb_release -sc) main" | sudo tee -a /etc/apt/sources.list.d/trivy.list
sudo apt-get update
sudo apt-get install trivy -y

```

```
#install maven
sudo apt-get update
sudo apt install openjdk-17-jdk -y
sudo apt install maven -y
```

## Appendix II

### Containerization & Kubernetes Deployment Manifests

#### II.1 Kubernetes Deployment Manifest (Kubernetec/regapp-deploy.yml)

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: regapp-deployment
  labels:
    app: regapp

spec:
  replicas: 2
  selector:
    matchLabels:
      app: regapp

  template:
    metadata:
      labels:
        app: regapp
    spec:
      containers:
        - name: regapp
          image: maxain27/register-app-pipeline
          imagePullPolicy: Always
          ports:
            - containerPort: 8080
      strategy:
        type: RollingUpdate
        rollingUpdate:
          maxSurge: 1
          maxUnavailable: 1
```

#### II.2 Kubernetes Service Manifest (Kubernetec/regapp-service.yml)

```
apiVersion: v1
kind: Service
metadata:
  name: regapp-service
  labels:
    app: regapp
spec:
  selector:
    app: regapp

  ports:
    - port: 8080
```

```
targetPort: 8080  
type: LoadBalancer
```

## II.3 Application Dockerfile

```
FROM tomcat:latest  
  
RUN cp -R /usr/local/tomcat/webapps.dist/* /usr/local/tomcat/webapps  
  
COPY webapp/target/*.war /usr/local/tomcat/webapps/
```